

IT workshop v Plzni

Co je nového v MySQL®



1. 2. 2023 v 18:00

Restaurace U Salzmannů

Salonek v 1. patře

Vstup a občerstvení zdarma



PeoplePath


Focus of this talk

- Interesting features from MySQL 5.7 and 8.0.
- Mainly from developer perspective.
- Some features will contain an example or possible use cases.

MySQL versioning

- Since 8.0 release quarterly.
- GA version since 8.0.11 (2018-04-19),
- latest version, 8.0.32 (2023-01-17).
- Does NOT FOLLOW semantic versioning.

Semantic versioning



MySQL versioning

- Does NOT FOLLOW semantic versioning.
- Breaking changes even in patch releases:
 - Removal of TLSv1 and TLSv1.1 in 8.0.28
 - MySQL Protocol changes in 8.0.24

But also a lot of new exciting features!

Exciting features

- Generated columns,
- JSON support,
- instant DDL,
- various index improvements,
- common table expression and
- window functions.

Generated columns

Generated columns

- since 5.7.5 (5.7.9 GA)
- Allows to store automatically generated data in a table.
- Column value is computed by predefined expression.
- Value cannot be changed manually and can be indexed.
- `GENERATED ALWAYS AS (expression) [STORED|VIRTUAL]`

Generated columns

```
CREATE TABLE users (  
    id          INT AUTO_INCREMENT PRIMARY KEY,  
    name        VARCHAR(60) NOT NULL,  
    surname     VARCHAR(60) NOT NULL,  
    full_name   VARCHAR(120)  
    GENERATED ALWAYS AS (CONCAT(name, ' ', surname))  
);
```

Generated columns

```
SELECT * FROM users;
```

id	name	surname	full_name
1	Jane	Doe	Jane Doe
2	Janie	Stiles	Janie Stiles
3	Richard	Miles	Richard Miles

Generated columns - few rules

- Used expression cannot:
 - reference another generated column,
 - use columns outside table,
 - contains non-deterministic functions (eg.: NOW()).

Generated columns - types

- **VIRTUAL** - evaluates its value on the fly (default).
- **STORED** - evaluates its value and store it on the disk.

```
CREATE TABLE users_alter (  
  id          INT AUTO_INCREMENT PRIMARY KEY,  
  name        VARCHAR(60) NOT NULL,  
  surname     VARCHAR(60) NOT NULL,  
  full_name   VARCHAR(120)  
  GENERATED ALWAYS AS (CONCAT(name , ' ', surname )) VIRTUAL,  
  hash varchar(32)  
  GENERATED ALWAYS AS (MD5(CONCAT(name , ' ', surname ))) STORED  
);
```

Generated columns - types

```
SELECT * FROM users_alter;
```

id	name	surname	full_name	hash
1	Jane	Doe	Jane Doe	f001124...
2	Janie	Stiles	Janie Stiles	55bac62...
3	Richard	Miles	Richard Miles	f32f3cd...

Generated columns - virtual

- ✅ They do not require disk space.
- ✅ INSERT and UPDATE queries come with no overhead.
- ❌ MySQL has to evaluate them during read operations.

Generated columns - stored

- ✓ No performance penalty during SELECT.
- ✗ INSERT or UPDATE comes with an overhead.
- ✗ They require disk space.

Use cases

- To simplify and unify queries.
- To cache a complicated conditions.
- To index a complex value.
- To extract a value from JSON data column.

JSON support

JSON support

- Since 5.7.9, more feature complete from 8.0
- Native JSON column data type.
- Cannot be indexed directly (functional indexes are way).
- Query language JSONPath.

JSONPath

- dot notation:

```
$.tool.jsonpath.creator.location[2]
```

- bracket notation:

```
$value['tool']['jsonpath']['creator']['location'][2]
```

JSON data type, advantages

- Advantages over text column:
 - automatic validation,
 - optimized storage format,
 - supports common operations (where condition, etc),
 - since 8.0.2 (8.0.11 GA) in-place update.

JSON data type

```
CREATE TABLE activity (  
    id          int auto_increment primary key,  
    event_name  ENUM('page-view', 'user-login'),  
    user_id     int,  
    properties  json,  
    browser     json  
);
```

JSON data type

```
INSERT INTO activity(event_name, user_id, properties, browser)
VALUES
(
    'page-view',
    1,
    '{ "page": 1 }',
    '{ "name": "Safari", "os": "Mac", "resolution": { "x": 1920, "y": 1080 } }'
),
(
    'page-view',
    2,
    '{ "page": 2 }',
    '{ "name": "Firefox", "os": "Windows", "resolution": { "x": 2560, "y": 1600 } }'
);
```

JSON data type, path operator

- column path operator (->) - shortcut to JSON_EXTRACT

```
SELECT id, browser->'$.name' browser  
FROM activity;
```

id	browser
1	"Safari"
2	"Firefox"

JSON data type, inline path operator

- column inline path operator (->>), since 8.0

```
SELECT id, browser->>'$.name' browser  
FROM activity;
```

id	browser	
1	Safari	
2	Firefox	

JSON data type, inline path operator, where

- inline path operator (->>) in where condition

```
SELECT id, browser->>'$.os' os
FROM activity
WHERE browser->>'$.name'='Firefox';
```

id	os
2	Windows

Partial update

- Since 8.0.2 (8.0.11 GA), `JSON_SET(column, path, value)`

```
UPDATE activity
SET `browser` = JSON_SET(
    `browser`,
    '$.name',
    'Phoenix'
)
WHERE browser->>'$.name'='Firefox';
```

[2022-10-02 22:18:41] 1 row affected in 4 ms

Partial update - rest of functions

- `JSON_SET()` replaces existing values and adds non existing values.
- `JSON_INSERT()` inserts values without replacing values.
- `JSON_REPLACE()` replaces only existing values.
- `JSON_REMOVE()` removes data from a JSON document.

Generated column + JSON = 🥰

```
CREATE TABLE activity (  
    id          int auto_increment primary key,  
    event_name  ENUM('page-view', 'user-login'),  
    user_id     int,  
    properties  json,  
    browser     json,  
    browser_name varchar(20)  
    GENERATED ALWAYS AS (`browser` ->> '$.name')  
);
```

Use cases

- Allows to mix document database with relation database.
- This can be a tricky!
- Possible use cases:
 - Error logging,
 - application event logging and
 - piloting ideas.

Instant DDL

Instant “Data definition language” - schema changing command

Instant DDL

- Partial support since 8.0.12, extended in 8.0.29
- Allows schema changes without making data unavailable.
- No need to do anything special to enable online DDL.

But you have to understand what is happening under the hood!

DDL Algorithms

- Algorithms InnoDB supports:
 - COPY
 - INPLACE
 - INSTANT

COPY Algorithm

- MySQL internally:
 - Create a new table with the altered schema.
 - Migrate data into new table.
 - Swaps the table names.
 - Drops the old table.

COPY Algorithm drawbacks

- Rollback of operation can be an expensive process.
- Concurrent DML's are not allowed during the ALTER table.
- Causes replication lag.

INPLACE Algorithm

- Operations are done in-place in the original table.
- Uses a temporary log file to track data changes by DML queries during the change.
- After in-place operation finishes the log will be applied.

INPLACE Algorithm drawbacks

- Large number of concurrent DML's can fail.
- Rollback of operation can be an expensive process.
- Causes replication lag.

INSTANT Algorithm

Performs only metadata changes => Doesn't touch the data file of the table.



...is not supported for all DDL operations.

INSTANT Algorithm, support

- Adding a column
- Dropping a column
- Renaming a column
- Modifying a column default value
- Renaming table

INSTANT Algorithm drawbacks

- Best works on latest patch version of MySQL, e.g:
 - Prior to 8.0.29, a column can only be added as a last one.
- You can do only 64 INSTANT operations on one table.

INSTANT Algorithm drawbacks

- Is not supported for:
 - tables that use ROW_FORMAT=COMPRESSED,
 - tables with a FULLTEXT index,
 - temporary tables and
 - stored columns.

INSTANT Algorithm recommendations

- Can be forced by `ALGORITHM=INSTANT;`
- Check patch version of MySQL:
 - at least 8.0.29 is recommended.
- Be aware of combinations:
 - column drop is instant, **index drop** not!

INSTANT Algorithm recommendations

- Always consult with documentation
- Part: “*15.12.1 Online DDL Operations*”:

<https://dev.mysql.com/doc/refman/8.0/en/innodb-online-ddl-operations.html>

Indexes

Indexes

- In MySQL 8, new types of index:
 - Multi-Valued,
 - functional,
 - descending and
 - invisible.

Multi-Valued Index

- Since 8.0.17
- Index defined on a column that stores an array of values.
- Use case - index JSON arrays.
- `CAST(... AS ... ARRAY)` in the index definition.

Multi-Valued Indexes, creation

```
CREATE TABLE customers (  
  id          BIGINT NOT NULL AUTO_INCREMENT PRIMARY KEY,  
  modified    DATETIME DEFAULT CURRENT_TIMESTAMP,  
  custinfo    JSON,  
  INDEX zips((CAST(custinfo->'$.zipcode' AS UNSIGNED ARRAY)))  
);
```

Multi-Valued Indexes, usage

- Condition functions:
 - MEMBER OF(json_array)
 - JSON_CONTAINS(target, candidate[, path])
 - JSON_OVERLAPS(json_doc1, json_doc2)

Multi-Valued Indexes, MEMBER OF

Succeed whether a given value is an element of json_array.

```
SELECT * FROM customers
WHERE 123 MEMBER OF(custinfo->'$.zipcode');
```

id	modified	custinfo
2	2019-06-29 22:23:12	{"user": "Jill", "user_id": 22, "zipcode": [94568, 123, 94582]}
3	2019-06-29 22:23:12	{"user": "Bob", "user_id": 31, "zipcode": [94477, 123]}
5	2019-06-29 22:23:12	{"user": "Ted", "user_id": 56, "zipcode": [123, 94582]}

Multi-Valued Indexes, JSON_CONTAINS

Succeeds whether a given JSON document is contained within a target JSON document.

```
SELECT * FROM customers WHERE  
JSON_CONTAINS(custinfo->'$.zipcode', CAST('[123,456]' AS JSON));
```

id	modified	custinfo
2	2019-06-29 22:23:12	{"user": "Jill", "user_id": 22, "zipcode": [94568, 123, 456]}
5	2019-06-29 22:23:12	{"user": "Ted", "user_id": 56, "zipcode": [123, 456]}

Multi-Valued Indexes, JSON_OVERLAPS

Succeed if the two document have a common array elements.

```
SELECT * FROM customers WHERE  
JSON_OVERLAPS(custinfo->'$.zipcode', CAST('[123,456]' AS JSON));
```

id	modified	custinfo
1	2019-06-29 22:23:12	{"user": "Jack", "user_id": 37, "zipcode": [456, 94536]}
2	2019-06-29 22:23:12	{"user": "Jill", "user_id": 22, "zipcode": [94568, 123, 456]}
3	2019-06-29 22:23:12	{"user": "Bob", "user_id": 31, "zipcode": [94477, 123]}
5	2019-06-29 22:23:12	{"user": "Ted", "user_id": 56, "zipcode": [123, 456]}

Functional indexes

- Since 8.0.13
- Function can be used as base for index.
- Usage - a filter condition against functional expression.

```
SELECT AVG(price) FROM products WHERE MONTH(create_time)=10;
```

Functional indexes

```
CREATE TABLE `products` (  
  `id`          int unsigned NOT NULL PRIMARY KEY AUTO_INCREMENT,  
  `price`       integer DEFAULT NULL,  
  `create_time` timestamp NULL DEFAULT NULL,  
  KEY `functional_index` ((month(`create_time`)))  
) ENGINE=InnoDB;
```

Functional indexes

```
EXPLAIN SELECT AVG(price) FROM products WHERE MONTH(create_time)=10;
```

without index

```
id: 1
select_type: SIMPLE
table: products
partitions: NULL
type: ALL
possible_keys: NULL
key: NULL
key_len: NULL
ref: NULL
rows: 1
filtered: 100.00
Extra: Using where
```

with index

```
id: 1
select_type: SIMPLE
table: products
partitions: NULL
type: ref
possible_keys: functional_index
key: functional_index
key_len: 5
ref: const
rows: 1
filtered: 100.00
Extra: NULL
```

Functional indexes, implementation

- They are implemented as hidden virtual generated columns:
 - They counts against the total limit of table columns.
 - Use only functions permitted for generated columns.
 - Subqueries, parameters, variables, stored functions, and loadable functions are not permitted.

Functional indexes + JSON = 💖

- `JSON_VALUE()` introduced in MySQL 8.0.21 allows to transparently index properties.

```
SELECT JSON_VALUE('{ "fname": "Joe", "lname": "Palmer" }', '$.fname');
```

```
+-----+
| JSON_VALUE('{ "fname": "Joe", "lname": "Palmer"}', '$.fname') |
+-----+
| Joe |
+-----+
```

Functional indexes + JSON = 🧡

```
SELECT JSON_VALUE(json_doc, path RETURNING type);
```

Equivalent to:

```
SELECT CAST(  
    JSON_UNQUOTE(JSON_EXTRACT(json_doc, path))  
    AS type  
);
```


Functional indexes + JSON = 💖

```
CREATE TABLE data(  
  j JSON,  
  INDEX i1 ( (JSON_VALUE(j, '$.id' RETURNING UNSIGNED)) )  
);
```

```
EXPLAIN SELECT * FROM data WHERE JSON_VALUE(j, '$.id' RETURNING UNSIGNED) = 123;
```

```
      id: 1  
select_type: SIMPLE  
      table: data  
  partitions: NULL  
        type: ref  
possible_keys: i1  
          key: i1  
      key_len: 9  
         ref: const  
        rows: 1  
   filtered: 100.00  
      Extra: NULL
```

Descending Indexes

- Since 8.0.1, (8.0.11)
- Stores key values in descending order.
- Can be combined in multiple-column indexes.
- Can increase the performance of following pattern:

```
ORDER BY field1 DESC, field2 ASC LIMIT N;
```

Descending Indexes

```
CREATE TABLE `articles` (  
  `id` int(11) NOT NULL AUTO_INCREMENT,  
  `name` varchar(100) DEFAULT NULL,  
  `created` datetime DEFAULT NULL,  
  PRIMARY KEY (`id`),  
  KEY `created_desc_name_asc` (`created` DESC, `name`)  
) ENGINE=InnoDB;
```

```
SELECT * FROM articles ORDER BY created DESC, name ASC limit 10;
```

id	name	created
1	foo	2022-10-01 16:20:52
3	quz	2022-06-18 16:21:27
2	bar	2022-06-09 16:21:08

Invisible indexes

- Maintained indexes that are not used by the optimizer.
- Allows to test the effect of removing an index, without making a destructive change.

```
ALTER TABLE t1 ALTER INDEX i_idx INVISIBLE;
```

```
ALTER TABLE t1 ALTER INDEX i_idx VISIBLE;
```

Indexes

- Other improvements:
 - histograms,
 - simultaneous index build (8.0.27) and
 - CHECK Constraints 🥰.

Intersect and except

Intersect and except

- since 8.0.31
- **INTERSECT** limits the result from multiple **SELECT** statements to those rows which are common to all.
- **EXCEPT** limits the result from the first **SELECT** statement to those rows which are (also) not found in the second.

Intersect and except

- As with UNION, the operands must have the same number of columns.
- DISTINCT modifier can remove duplicates from either side of the intersection.

Intersect and except

```
SELECT * FROM ordered_food;
```

id	name	tacos	sushis
1	Kenny	NULL	10
2	Miguel	5	0
3	Bohus	4	5
4	Kajiyamasan	NULL	10
5	Scott	10	NULL
6	Lenka	NULL	NULL

Intersect

```
SELECT * FROM ordered_food  
WHERE tacos > 0;
```

id	name	tacos	sushis
2	Miguel	5	0
3	Bohus	4	5
5	Scott	10	NULL

```
SELECT * FROM ordered_food  
WHERE sushis > 0;
```

id	name	tacos	sushis
1	Kenny	NULL	10
3	Bohus	4	5
4	Kajiyamasan	NULL	10

Intersect

```
SELECT * FROM ordered_food WHERE tacos>0  
INTERSECT  
SELECT * FROM ordered_food WHERE sushis>0;
```

id	name	tacos	sushis
3	Bohus	4	5

Except

- limits the result from the first statement to rows which are not found in the second.

```
SELECT * FROM ordered_food  
WHERE tacos > 0;
```

id	name	tacos	sushis
2	Miguel	5	0
3	Bohus	4	5
5	Scott	10	NULL

```
SELECT * FROM ordered_food  
WHERE sushis > 0;
```

id	name	tacos	sushis
1	Kenny	NULL	10
3	Bohus	4	5
4	Kajiyamasan	NULL	10

Except

```
SELECT * FROM ordered_food WHERE tacos>0  
EXCEPT  
SELECT * FROM ordered_food WHERE sushis>0;
```

id	name	tacos	sushis
2	Miguel	5	0
5	Scott	10	NULL

CTE - Common Table Expression

CTE - Common Table Expression

- since 8.0.1, (8.0.11)
- An alternative to a derived table and view.
- Simplifies complex joins and subqueries.
- syntax: `WITH <name> AS (<query>)`


CTE, motivation

- Better readability of the queries.
- Improved performance.
- A valid alternative to a VIEW and a temporary table.
- Possibility to create recursive queries.

CTE, example with EAV

- Entity–attribute–value model

id	name	surname
1	John	Doe



Entity (user_id)	Attribute	Value
1	name	John
1	surname	Doe

CTE, example with EAV

- Wordpress eg.: wp_usermeta
- Transform from EAV to regular table can be done by pivoting.

umeta_id	user_id	meta_key	meta_value
1	1	first_name	Emma
2	1	last_name	Obrien
3	1	nickname	admin

CTE, example with EAV

```
SELECT user_id,  
       MAX(CASE WHEN meta_key='first_name' THEN meta_value END) as first,  
       MAX(CASE WHEN meta_key='last_name' THEN meta_value END) as last  
FROM `wp_usermeta`  
GROUP BY user_id
```

user_id	first	last
1	Emma	Obrien
2	Nial	Casey
3	Keeley	Brookes
4	Bert	Mccoy
5	Alyce	Sheldon

CTE, example with EAV

- CTE with pivot query

```
WITH cte AS (  
    SELECT user_id,  
           MAX(CASE WHEN meta_key='first_name' THEN meta_value END) as first,  
           MAX(CASE WHEN meta_key='last_name' THEN meta_value END) as last  
    FROM `wp_usermeta`  
    GROUP BY user_id  
)  
SELECT * FROM cte;
```

CTE, example with EAV

- CTE with pivot query

```
WITH cte AS (  
    SELECT user_id,  
           MAX(CASE WHEN meta_key='first_name' THEN meta_value END) as first,  
           MAX(CASE WHEN meta_key='last_name' THEN meta_value END) as last  
    FROM `wp_usermeta`  
    GROUP BY user_id  
)  
SELECT * FROM cte  
WHERE first = "Emma" OR last = "Sheldon"  
ORDER BY first
```

CTE, recursive

- A recursive CTE with subquery that refers itself.
- Usage:
 - to generate series and
 - hierarchical or tree-structured data traversal.

CTE, recursive

- A recursive CTE components:

```
WITH RECURSIVE cte AS (  
    initial_query      -- "seed" member  
    UNION ALL  
    recursive_query    -- recursive member referring the same CTE  
)  
SELECT * FROM cte;    -- main query
```

CTE, recursive

- **Seed member** - initial query, executed in first iteration.
- **Recursive member** - contains the reference to the same CTE name.
- This second component will generate all the remaining items of the main query.

CTE, recursive, example

```
WITH RECURSIVE cte (n) AS
```

```
(
```

```
  SELECT 1
```

```
  UNION ALL
```



```
  SELECT n + 1 FROM cte WHERE n < 5
```

```
)
```

```
SELECT * FROM cte;
```

-- "seed" member

-- recursive member

-- main query

n
1
2
3
4
5

CTE, recursive, example

```
WITH RECURSIVE cte (n) AS
(
    SELECT '2013-01-01'
    UNION ALL
    SELECT n + INTERVAL 1 DAY FROM cte WHERE n < '2013-01-10'
)
SELECT * FROM cte;
```

n
2013-01-01
2013-01-02
2013-01-03
2013-01-04
...

CTE, recursive, example

```
CREATE TABLE orgchart(  
    id          INT PRIMARY KEY AUTO_INCREMENT,  
    name        VARCHAR(20),  
    role        VARCHAR(20),  
    manager_id  INT,  
    FOREIGN KEY (manager_id) REFERENCES orgchart(id)  
);
```

CTE, recursive, example

id	name	role	manager_id
1	Matthew	CEO	NULL
2	Caroline	CFO	1
3	Tom	CTO	1
4	Sam	Treasurer	2
5	Ann	Controller	2
6	Anthony	Dev Director	3
7	Lousie	Sys Admin	3
8	Travis	Senior DBA	3
9	John	Developer	6
10	Jennifer	Developer	6
11	Maria	Junior DBA	8

CTE, recursive, example

id	name	path	level
1	Matthew	Matthew	1
2	Caroline	Matthew -> Caroline	2
3	Tom	Matthew -> Tom	2
4	Sam	Matthew -> Caroline -> Sam	3
5	Ann	Matthew -> Caroline -> Ann	3
6	Anthony	Matthew -> Tom -> Anthony	3
7	Lousie	Matthew -> Tom -> Lousie	3
8	Travis	Matthew -> Tom -> Travis	3
9	John	Matthew -> Tom -> Anthony -> John	4
10	Jennifer	Matthew -> Tom -> Anthony -> Jennifer	4
11	Maria	Matthew -> Tom -> Travis -> Maria	4

CTE, recursive, example

```
WITH RECURSIVE reporting_chain(id, name, path, level) AS (  
    SELECT id, name, CAST(name AS CHAR(100)), 1  
    FROM orgchart  
    WHERE manager_id IS NULL  
    UNION ALL  
    SELECT oc.id, oc.name, CONCAT(rc.path, ' -> ', oc.name), rc.level+1  
    FROM reporting_chain rc JOIN orgchart oc ON rc.id=oc.manager_id)  
SELECT * FROM reporting_chain ORDER BY level;
```

CTE, recursive, example

id	name	path	level
1	Matthew	Matthew	1
2	Caroline	Matthew -> Caroline	2
3	Tom	Matthew -> Tom	2
4	Sam	Matthew -> Caroline -> Sam	3
5	Ann	Matthew -> Caroline -> Ann	3
6	Anthony	Matthew -> Tom -> Anthony	3
7	Lousie	Matthew -> Tom -> Lousie	3
8	Travis	Matthew -> Tom -> Travis	3
9	John	Matthew -> Tom -> Anthony -> John	4
10	Jennifer	Matthew -> Tom -> Anthony -> Jennifer	4
11	Maria	Matthew -> Tom -> Travis -> Maria	4

CTE, recursive, limitation

- **cte_max_recursion_depth** - max recursion depth (1000).
- **max_execution_time** - execution timeout for SELECT statements.
- **MAX_EXECUTION_TIME** - optimizer hint enforces a per-query.

Window functions

Window functions

- since 8.0.2, (8.0.11)
- Offers aggregate-like functionality on a defined range of rows in a query.
- Window functions will return a value for every row in a query result.

Window functions

- **OVER** - indicates usage of window function
- **PARTITION BY** - marks how to divide the rows into groups

```
SELECT
    <agregation>(field) OVER() AS field_name,
    <agregation>(field) OVER(PARTITION BY field) AS field_name
FROM <table name>
```

Window functions, example

```
CREATE TABLE sales(  
    id          INT PRIMARY KEY AUTO_INCREMENT,  
    year        INT,  
    country     VARCHAR(20),  
    product     VARCHAR(32),  
    profit      INT  
);
```

Window functions, example

```
SELECT * FROM sales ORDER BY country, year, product;
```

id	year	country	product	profit
1	2000	Finland	Computer	1500
2	2000	Finland	Phone	100
3	2001	Finland	Phone	10
4	2000	India	Calculator	75
5	2000	India	Calculator	75
6	2000	India	Computer	1200
7	2000	USA	Calculator	75
8	2000	USA	Computer	1500
9	2001	USA	Calculator	50
10	2001	USA	Computer	1500
11	2001	USA	Computer	1200
12	2001	USA	TV	150
13	2001	USA	TV	100

Window functions, regular aggregations

```
SELECT SUM(profit) AS total_profit FROM sales;
```

total_profit
7535

```
SELECT country, SUM(profit) AS country_profit  
FROM sales GROUP BY country ORDER BY country;
```

country	country_profit
Finland	1610
India	1350
USA	4575

Window functions, example

```
SELECT
    year, country, product, profit,
    SUM(profit) OVER() AS total_profit,
    SUM(profit) OVER(PARTITION BY country) AS country_profit
FROM sales
ORDER BY country, year, product, profit;
```

Window functions, example

year	country	product	profit	total_profit	country_profit
2000	Finland	Computer	1500	7535	1610
2000	Finland	Phone	100	7535	1610
2001	Finland	Phone	10	7535	1610
2000	India	Calculator	75	7535	1350
2000	India	Calculator	75	7535	1350
2000	India	Computer	1200	7535	1350
2000	USA	Calculator	75	7535	4575
2000	USA	Computer	1500	7535	4575
2001	USA	Calculator	50	7535	4575
2001	USA	Computer	1200	7535	4575
2001	USA	Computer	1500	7535	4575
2001	USA	TV	100	7535	4575
2001	USA	TV	150	7535	4575

Window functions, example

```
SELECT
    year, country, product, profit,
    RANK() OVER(ORDER BY `profit` desc, `id`) total_profit_rank
FROM sales
ORDER BY total_profit_rank;
```

year	country	product	profit	total_profit_rank
2000	Finland	Computer	1500	1
2000	USA	Computer	1500	2
2001	USA	Computer	1500	3
2000	India	Computer	1200	4
2001	USA	Computer	1200	5
2001	USA	TV	150	6
2000	Finland	Phone	100	7

Window functions, example

```
SELECT
    year, country, product, profit,
    RANK() OVER(PARTITION BY `country` ORDER BY `profit` desc, id)
    total_profit_rank
FROM sales
ORDER BY country, total_profit_rank;
```

Window functions, example

year	country	product	profit	total_profit_rank
2000	Finland	Computer	1500	1
2000	Finland	Phone	100	2
2001	Finland	Phone	10	3
2000	India	Computer	1200	1
2000	India	Calculator	75	2
2000	India	Calculator	75	3
2000	USA	Computer	1500	1
2001	USA	Computer	1500	2
2001	USA	Computer	1200	3
2001	USA	TV	150	4
2001	USA	TV	100	5
2000	USA	Calculator	75	6
2001	USA	Calculator	50	7

Window functions, example

```
SELECT
```

```
    year, country, product, profit,
```

```
    RANK() OVER(
```

```
        PARTITION BY `country`
```

```
        ORDER BY `profit` desc, id
```

```
    ) total_profit_rank,
```

```
    profit - FIRST_VALUE( profit ) OVER (
```

```
        PARTITION BY `country`
```

```
        ORDER BY `profit` desc, id
```

```
    ) profit_back_of_first
```

```
FROM sales
```

```
ORDER BY country, total_profit_rank;
```

Window functions, example

year	country	product	profit	total_profit_rank	profit_back_of_first
2000	Finland	Computer	1500	1	0
2000	Finland	Phone	100	2	-1400
2001	Finland	Phone	10	3	-1490
2000	India	Computer	1200	1	0
2000	India	Calculator	75	2	-1125
2000	India	Calculator	75	3	-1125
2000	USA	Computer	1500	1	0
2001	USA	Computer	1500	2	0
2001	USA	Computer	1200	3	-300
2001	USA	TV	150	4	-1350
2001	USA	TV	100	5	-1400
2000	USA	Calculator	75	6	-1425
2001	USA	Calculator	50	7	-1450

Window functions, functions

CUME_DIST()

NTH_VALUE()

DENSE_RANK()

NTILE()

LEAD()

PERCENT_RANK()

LAG()

RANK()

LAST_VALUE()

ROW_NUMBER()

Other features

(also interesting...)

Other features

- Explain analyze,
- spatial Data Types,
- invisible columns and
- vastly improved replication.

Overview

Overview

- Even patch version can contains interesting features.
- Check which patch version you are running.
- You should consult documentation and release notes.



Interesting sources

- MySQL release notes -

<https://dev.mysql.com/doc/relnotes/mysql/8.0/en/preface.html>

- Percona blog -

<https://www.percona.com/blog/>



Questions?

Slides:

<http://bohuslav.simek.si/pph-mysql-new-features/>



Informační systémy
na míru



Business
analytika



Mobilní
aplikace



Weby
a E-shopy

Jak na čistý kód a škálovatelnou architekturu kódu

1. března 2023
18:00 hod



Bonus slides

Maria DB compatibility

Feature	Status	Version	Note
Generated columns	✓	10.2.	syntax compatible
JSON support	✓	10.2.7	syntax compatible, basic functionality
instant DDL	⚠	10.4	to some degree
common table expression	✓	10.2.1	
window functions	✓	10.2	

Maria DB compatibility

Feature	Status	Version	Note
Intersect, except	✓	10.3	Maria DB has a longer support
Multi-Valued index	✗		
Functional index	✗		
Descending index	✗		
Invisible index	✓	10.6	Ignored indexes is pretty much same